
Stakeholder Deliverable

Client 3: Final Report

Stakeholder-facing summary of the refreshed PSID crisis-module package

Document status

Prepared for client review

Updated: April 2026

Repository: Team-PSID

Contents

1 Executive Summary 1

2 Headline Results 1

3 Figures and Visual Diagnostics 2

4 Reporting Pipeline 4

5 Methodology in Plain Language 4

5.1 How ranking works 4

5.2 How the binary category is assigned 4

6 Threshold Diagnostics 5

7 Selected Module Composition 6

8 Why Financial Questions Do Not Dominate the Final Module 6

9 Demographic Note 7

10 Worked Example of a High-Value Generic Item 7

11 Deliverables for Client Use 8

12 Conclusion 8

1 Executive Summary

This final report explains the refreshed PSID crisis-module package in stakeholder terms. The package now uses a simpler category model, a cleaner deployable questionnaire, and a more defensible explanation of why some items survive the final shortlist.

Headline takeaway

The current production run starts from 52 historical questions and produces a 28-question scored module. The client-facing category model has been simplified to two labels only: **Generic** and **Specific**. Financial items now compete inside the broader **Specific** bank alongside disaster and pandemic items rather than appearing as a stand-alone final category.

2 Headline Results

Metric	Value
Total ranked questions	52
Selected scored questions	28
Deployable questions	26
Selected scored minutes	29.98
Deployable minutes	28.82
Generic rows	3
Specific rows	49
Average R_i	2.201
Maximum R_i	5.667
Average P_i	2.561
Maximum P_i	7.860

Table 1: Top-line performance metrics for the refreshed package.

3 Figures and Visual Diagnostics




Figure 1: Binary Generic versus Specific distribution.




Figure 2: Time-budget allocation across the selected module.




Figure 3: Utility-versus-burden relationship for the refreshed corpus.




Figure 4: Top-ranked items retained in the final package.

4 Reporting Pipeline




Figure 5: Reporting flow from historical question bank to the current client bundle.

5 Methodology in Plain Language

5.1 How ranking works

Each question receives a utility score for crisis relevance, a burden score for length and complexity, a baseline efficiency score R_i , and an augmented priority score P_i that rewards distinctiveness and construct coverage.

Main formulae

$$R_i = \frac{U_i}{B_i}$$

$$B_i = \max(0.10 \times \text{word_count} + 0.20 \times \text{complexity}, 0.1)$$

$$P_i = \frac{\text{augmented_utility}}{B_i \times (1 + 0.65 \times \text{redundancy_scaled})}$$

5.2 How the binary category is assigned

The client requested a simple threshold on R_i , but raw R_i already exceeds 0.70 for the entire refreshed corpus. The workflow therefore applies the cutoff to normalized R_i instead.

Operational thresholding

$$\text{ri_threshold_score} = \frac{R_i - \min(R_i)}{\max(R_i) - \min(R_i)}$$

Generic $\iff \text{ri_threshold_score} \geq 0.70$

Specific $\iff \text{ri_threshold_score} < 0.70$

6 Threshold Diagnostics

Diagnostic	Value
Minimum raw R_i	0.818
Maximum raw R_i	5.667
Rows with raw $R_i \geq 0.70$	52 of 52
Rows with normalized threshold score ≥ 0.70	3 of 52

Table 2: Why the Generic bucket is intentionally small after normalization.

This is why the current **Generic** bucket contains only the most portable high-signal items after normalization.

7 Selected Module Composition

Final category	Source	Selected questions	Minutes	Average R_i	Average P_i
Generic	COVID-19	3	1.40	5.333	7.125
Specific	Hurricane Katrina 2007	14	20.30	2.617	2.769
Specific	COVID-19	5	2.92	3.364	4.393
Specific	Govt Shutdown Income	3	2.33	3.326	4.482
Specific	Govt Shutdown Crisis	1	1.87	3.719	4.592
Specific	Understanding Society	2	1.17	1.447	2.329

Table 3: Composition of the current scored module by final category and historical source.

8 Why Financial Questions Do Not Dominate the Final Module

The earlier output looked more finance-heavy than the current one for two reasons. First, **Financial Crisis** used to appear as its own visible final label. Second, multiple finance-related items were close variants of one another, which made the previous structure look broader than it really was.

Interpretation

- Financial items now compete within **Specific** instead of standing alone as a final category.
- Redundancy penalties reduce the advantage of repeating near-identical income questions.
- Katrina disaster items contribute many strong direct-exposure and trauma items.
- The three highest-portability Generic items still capture crisis-related financial disruption without making finance the whole story.

9 Demographic Note

Group	Rows	Average R_i	Average P_i
Demographics only	3	1.354	2.189
Non-demographic crisis items	49	2.253	2.584

Table 4: Demographic rows versus direct crisis-impact rows.

- Demographic rows provide context and traceability.
- They are not as efficient as direct crisis-impact questions for the deployable instrument.
- That is why the final scored module can retain them while the deployable questionnaire removes them.

10 Worked Example of a High-Value Generic Item

Example question: GEN_EXPERIENCED_FINANCIAL_DIFFICULTIES_C

Field	Value	Meaning
R_i	5.667	Very high baseline signal per unit of burden.
P_i	7.187	Strong augmented priority after distinctiveness adjustments.
minutes	0.35	Short enough to include in the universal core.
Final category	Generic	Portable across crisis types.

Table 5: Characteristics of a high-value Generic item.

This is the kind of item the new **Generic** bucket is meant to preserve: short, portable, and strongly informative.

11 Deliverables for Client Use

Deliverable	Use
docs/client 1.md	Technical Markdown codebook with formulas, figures, schema, and workflow notes.
docs/client 2.md	Markdown deployable questionnaire with live question tables and universe rules.
docs/client 3.md	Markdown final report for stakeholder circulation.
deliverables/Master_Questionnaire.xlsx	Final internal review sheet with traceability fields.
deliverables/Deployable_Questionnaire.pdf	Final deployable PDF questionnaire.
deliverables/Codebook.pdf	Rendered technical codebook PDF from the artifact pipeline.
deliverables/Final_Report.pdf	Rendered stakeholder-facing PDF from the artifact pipeline.
deliverables/Client_1_Technical_Codebook.pdf	LaTeX codebook of the technical codebook.
deliverables/Client_2_Deployable_Questionnaire.pdf	LaTeX questionnaire of the deployable questionnaire.
deliverables/Client_3_Final_Report.pdf	LaTeX report of this stakeholder report.

Table 6: Client-facing outputs available after the LaTeX build.

12 Conclusion

The refreshed package is more disciplined than the earlier version in three ways:

- It keeps the category model easy to explain.
- It preserves the quantitative ranking logic instead of replacing it with ad hoc pruning.
- It cleanly separates technical documentation, deployable wording, and stakeholder reporting.

The result is a client package that is easier to defend analytically and easier to use operationally.